

Label Noise and Training-Set Size in Small-U-Net Segmentation of Synthetic Shapes: Object Flips Hurt, Boundary Jitter Does Not, and Classification-Style Robust Losses Barely Help

Tern-81A

STEM seed corpus

<https://github.com/mihir-chauhan/research-harness-papers/tree/main/papers/label-noise-and-data-size-in-segmentation-on-syn>

Abstract—We ask how training-set size and the kind of label noise jointly determine the accuracy of a small U-Net on synthetic 64×64 images of overlapping circles, squares and triangles, where the clean ground truth is known exactly. Training masks are corrupted either by boundary jitter or by whole-object class flips (nominal level up to 30%); training-set sizes range from 100 to 2000 images; we compare plain cross-entropy (CE) with reimplemented GCE and SCE losses and a simple band-ignore loss. On clean labels mIoU grows from 0.68 at $N = 100$ to 0.82 at $N = 2000$ (Welch $p = 0.001$, four seeds). At $N = 500$, boundary jitter at level 0.3 costs CE only 0.01 mIoU, whereas object flips at the same level cost 0.27 (registered Welch test of flip vs boundary mIoU: $p = 0.011$). More data did not dilute flip noise in our runs: the flip-induced drop is not smaller at $N = 2000$ than at $N = 100$ (Welch $p = 0.25$, inconclusive), and a training-length sweep up to 1200 steps shows no late-training collapse under flips, although the network fits the flipped labels at $N = 100$. No robust loss is reliably better than CE: GCE has the higher mean mIoU under flip noise (CE minus GCE: -0.07), which is not significant with four seeds ($p = 0.38$), SCE is indistinguishable from CE, and our band-ignore loss is clearly worse on clean and boundary-noise labels. This is a small CPU study (four seeds in the main grid, three in the sweeps, one architecture, one synthetic domain), and per-seed variance under flip noise is large.

Index Terms—semantic segmentation, label noise, training-set size, noise-robust losses, synthetic data, U-Net

I. Introduction

Dense prediction needs dense labels, and dense labels are expensive and imperfect: annotators trace object outlines a few pixels off, and crowd-sourced or semi-automatic pipelines occasionally assign a whole region to the wrong class. Practitioners therefore face two coupled questions: how many (noisy) annotated images are needed, and does the kind of noise matter? Work on label noise has mostly studied image classification [1]–[4], where a network first

fits the clean pattern and only later memorises wrong labels, and where very large datasets can average out even massive uniform noise. Segmentation differs: a label error is a spatially structured set of pixels, boundary errors touch only a thin band around each object, and the effective number of training “samples” is the number of pixels, not of images.

We run a controlled, small-scale study of this interaction on a problem in which the ground truth is known exactly. We generate 64×64 images of overlapping circles, squares and triangles on textured backgrounds, corrupt only the training masks with two noise models (boundary jitter and whole-object class flips, each at a nominal level of up to 30%), vary the training-set size from 100 to 2000 images, and train a small U-Net [5] on a CPU. We measure per-class IoU against clean test masks. Besides plain cross-entropy (CE) we reimplement two noise-robust losses from the classification literature, generalized cross-entropy (GCE) [6] and symmetric cross-entropy (SCE) [7], and add a simple band-ignore CE that drops the pixels near annotated edges.

We register four hypotheses beforehand (data size helps; object flips hurt more than boundary jitter; noise hurts more with little data; robust losses help) and report each verdict with the test we committed to, including where the evidence is thin. Contributions: (i) a reproducible synthetic benchmark and noise generators; (ii) a size $\times$ noise-type $\times$ noise-level grid with 3–4 seeds per cell and per-class IoU; (iii) a head-to-head of CE, GCE, SCE and band-ignore CE under identical budgets; (iv) an honest account of what a CPU-scale study of this kind can and cannot establish.

II. Related Work

Segmentation backbones and losses. U-Net [5] was designed for segmentation from few annotated images with heavy augmentation, which makes it the natural model for a data-size study; we use a much smaller variant. Overlap-

Tern-81A is the codename of an autonomous study, not a person. Produced by the STEM autonomous AI research pipeline; the topic was chosen to cover the field. Every reported result is traced to logged runs in the linked repository and every reference was checked against arXiv, Crossref, DataCite or OpenAlex.

based losses such as the Dice loss of V-Net [8] are common for class imbalance; we keep pixel-wise losses so that the noise-robust variants are directly comparable, and leave Dice-type losses out of scope.

Memorisation and noisy labels. Deep networks can fit random labels [1], but on real data they tend to learn simple, clean patterns before memorising noise [2]; classification accuracy can remain high under massive uniform noise if the dataset is large enough [3]. This motivates our interaction between noise and data size and our training-length sweep. Surveys of the field are given by Song et al. [4].

Noise-robust training. GCE [6] interpolates between CE and the mean absolute error, SCE [7] adds a reverse cross-entropy term, Co-teaching [9] trains two networks that select small-loss samples for each other, and early-learning regularisation [10] exploits the early-learning phenomenon. We reimplement GCE and SCE (they are one-line pixel-wise losses and need no extra network); Co-teaching and ELR need sample-level selection or per-sample memory that does not map cleanly to pixel-wise noise in a small study and are not run.

Segmentation under noisy annotation and data scaling. Karimi et al. [11] survey techniques for learning with noisy labels in medical image analysis, including segmentation, which is the closest prior work to ours; our contribution is the controlled synthetic setting where noise type, level and training-set size are crossed with exactly known ground truth. Empirical data-scaling studies [12] motivate reading performance as a function of N .

III. Method

A. Problem and noise models

Let $x \in [0, 1]^{3 \times H \times W}$ be an image with $H = W = 64$ and $y \in \{0, \dots, 3\}^{H \times W}$ its clean label map (background, circle, square, triangle; the label of a pixel is that of the top-most visible object). Only the training labels are corrupted, $\tilde{y} = \mathcal{C}(y)$; test labels are always clean. We define two corruptions at nominal level p :

Boundary noise. Let $B(y) = \{u : \max_{v \in \mathcal{N}_5(u)} y_v \neq \min_{v \in \mathcal{N}_5(u)} y_v\}$ be the set of pixels within 2 px of a label edge ($\mathcal{N}_5$ is the 5×5 neighbourhood). Each $u \in B(y)$ independently, with probability p , receives the label of a uniformly chosen pixel of $\mathcal{N}_5(u)$; other pixels are unchanged. At the same p boundary noise changes a tiny fraction of pixels (reported as `label_noise_pixels`).

Object-flip noise. Each visible object (instance with at least one visible pixel) is, with probability p , assigned a uniformly random class different from its own for all its visible pixels. This is the object-level analogue of symmetric class noise.

B. Network and losses

All systems share one U-Net [5]: three down-sampling stages (max-pool) and a bottleneck with widths 8, 16, 32, 64, transposed-convolution up-sampling with skip

connections, batch normalisation, a 1×1 output convolution to four logits, and 120356 parameters. Let $p_k(u)$ be the softmax probability of class k at pixel u , $\tilde{y}_u$ the (noisy) training label and $p_{\tilde{y}}(u) = p_{\tilde{y}_u}(u)$. The loss is a mean over pixels $\ell = \frac{1}{\sum_u w_u} \sum_u w_u \ell_u$ with:

$$\ell_u^{\text{CE}} = -\log p_{\tilde{y}}(u), \quad w_u = 1, \quad (1)$$

$$\ell_u^{\text{GCE}} = \frac{1 - p_{\tilde{y}}(u)^q}{q}, \quad q = 0.7, \quad (2)$$

$$\ell_u^{\text{SCE}} = \alpha \ell_u^{\text{CE}} + \beta A (p_{\tilde{y}}(u) - 1), \quad A = -4, \quad (3)$$

$$\ell_u^{\text{BI}} = \ell_u^{\text{CE}}, \quad w_u = \mathbb{1}[u \notin B(\tilde{y})]. \quad (4)$$

GCE follows Zhang and Sabuncu [6]; it tends to CE as $q \rightarrow 0$ and to the mean absolute error at $q = 1$. SCE follows Wang et al. [7]: the reverse CE $-\sum_k p_k \log \tilde{q}_k$ with the one-hot target clipped to $\log 0 := A$ equals $-A(1 - p_{\tilde{y}})$, which is what Eq. (3) implements, with $\alpha = 0.1$ and $\beta = 1.0$. Both are reimplemented here, pixel-wise, with the defaults of the original papers. Band-ignore CE (BI) is our own simple method: it removes the loss on the band $B(\tilde{y})$ around the edges of the noisy mask, which is computed from the labels alone (no clean information), in the spirit of void/trimap labels. It cannot help against object flips, since the interior of a flipped object stays supervised, and it forfeits supervision of true edges.

C. Metrics

IoU of class k on the whole test set (pooled over pixels): $\text{IoU}_k = \frac{|\hat{Y}=k \wedge Y=k|}{|\hat{Y}=k \vee Y=k|}$; mIoU is the mean over the four classes (background included), `miou_fg` the mean over the three shapes. We also log pixel agreement of the final network with the noisy and with the clean training labels on the training images as a memorisation indicator.

IV. Experimental Setup

Data. Images are generated with numpy: a background of three smoothed-noise channels plus a sinusoidal grating, 3–5 objects of random class, size, position and orientation (circle; rotated square; equilateral triangle), each filled with a random colour plus its own smooth texture, pasted in random order so that objects overlap and occlude, and Gaussian pixel noise. A fixed pool of 2000 training images (generator seed 1000) and a fixed test set of 500 images (seed 2000) are used by every run. A run with size N draws N images from the pool with a permutation seeded by the run seed, so seed s at size N always sees the same subset and noise realisation across loss functions; noise is applied once, before training, with a seed that depends on the run seed.

Training. AdamW with peak learning rate 0.010 (OneCycle schedule), weight decay 0.0001, batch size 16, 600 optimiser steps for every N (so small N is trained for many more epochs than large N ; this fixed-compute choice is a design decision of the study and is probed in the training-length sweep), random horizontal flip, final checkpoint, no early stopping and no validation set (a

clean validation set would not exist in the noisy-label setting). All runs use two CPU threads on a shared machine; a run takes roughly one to two minutes.

Experiments. main: $N = 500$, four systems (CE, GCE, SCE, band-ignore CE) $\times$ three tasks (clean, boundary noise $p = 0.3$, object-flip noise $p = 0.3$), 4 seeds (0–3). Band-ignore CE additionally ran under flip noise $p = 0.1$ and $p = 0.2$ (three seeds), so that the matching CE sweep has a reference. size: CE only, $N \in \{100, 250, 1000, 2000\}$; clean and flip tasks at all sizes and boundary noise at $N = 100, 2000$, four seeds at $N = 100$ and $N = 2000$ and three seeds at $N = 250, 1000$ (the $N = 500$ column is the CE row of main). sweep_noise: CE, $N = 500$, flip $p \in \{0.1, 0.2\}$, three seeds. sweep_gce_q: GCE with $q \in \{0.3, 0.9\}$ under flip $p = 0.3$, $N = 500$, three seeds ($q = 0.7$ is the four-seed main row). sweep_steps: CE at $N = 100$ under flip $p = 0.3$ with 200 and 1200 steps, three seeds (the 600-step row is in size). Seeds fix the data subset, the noise, the initialisation and the batch order; every metric is reproducible from them. The seed count is below the five suggested in the task description because of the runtime budget; the study has no boundary-noise levels other than $p = 0.3$ and no noise-level sweep for boundary noise.

Pilots (disclosed). Two earlier pilots (300 steps at learning rate 0.005, then 450 steps at the same rate) left several seeds with a collapsed triangle or square class even on clean labels, so they could not show noise effects. Their rows are logged and marked superseded; the final budget (600 steps, learning rate 0.01) was chosen from them and from two single-seed clean-label runs that were scored on the test split. This is a small test-split leak in the choice of one global optimiser budget; it is the same for every system and noise condition, and no per-system or per-noise tuning was done. GCE and SCE use the defaults of the original papers.

Hardware and runtime. Shared Apple-silicon CPU, two threads per run, two runs at a time; one run takes roughly one to two minutes and the whole study (including pilots) about two hours of wall-clock time.

Statistics. Verdicts use Welch t -tests ($\alpha = 0.05$, no multiplicity correction) between per-seed values, as registered in proposal.md; for H4 the tests are those printed by rh compare against CE. With 3–5 seeds these tests have little power and we report standard deviations over seeds. All statistics come from logged runs.

V. Results

We go hypothesis by hypothesis. All tests are the Welch t -tests registered in proposal.md, with four seeds per arm and no multiplicity correction; p -values below are descriptive given the very small n .

H1 (more clean data helps): supported. CE on clean labels reaches mIoU 0.68 (± 0.04) at $N = 100$ and 0.82 (± 0.03) at $N = 2000$; the difference is 0.14 (Welch $p = 0.001$). The curve (Fig. 2, Table III) is concave: most of




Fig. 1. mIoU of the four systems at $N = 500$ (mean and spread over seeds; generated from the main group).




Fig. 2. CE test mIoU against training-set size (mean $\pm$ std over seeds). Boundary noise was run only at $N = 100, 500, 2000$.

the gain is reached by $N = 500$ (0.80), and $N = 1000$ gives 0.81. Square and triangle IoU rise most (Table III).

H2 (flips hurt more than boundary jitter at $N = 500$): supported. CE obtains 0.78 under boundary noise and 0.53 under flip noise at the same nominal level; the difference is 0.25 (Welch $p = 0.011$). Boundary noise is not distinguishable from clean labels at $N = 500$ (drop 0.01, Welch $p = 0.61$). The comparison is confounded by the fact that “the same nominal level” touches very different numbers of pixels (the fraction of training pixels whose label is changed, `label_noise_pixels`, is 0.018 for boundary noise and 0.076 for flips at $N = 500$), so the result says the two models at $p = 0.3$ differ, not that a flipped pixel is intrinsically more harmful than a jittered one. Fig. 3 shows that flips damage square and triangle most (CE square IoU 0.71 clean vs 0.28 under flips; triangle 0.67 vs 0.32), circle less (0.82 vs 0.54), and leave background IoU essentially unchanged (0.98 vs 0.98); the reported mIoU therefore hides that the damage is entirely a foreground-

TABLE I

Main results: test IoU (mean $\pm$ std over seeds) at $N = 500$. Bold: best, underline: second-best per task. Band-ignore CE at flip $p = 0.1, 0.2$ has three seeds; there is no CE/GCE/SCE row for those tasks in this group (CE is in Table IV).

Method	Task	miou $\uparrow$	iou_circle $\uparrow$	iou_square $\uparrow$	iou_triangle $\uparrow$	n
CE	none	0.80 ± 0.03	0.82 ± 0.03	0.71 ± 0.04	0.67 ± 0.07	4
SCE	none	0.77 ± 0.02	0.80 ± 0.02	0.69 ± 0.02	0.60 ± 0.08	4
GCE	none	0.79 ± 0.02	0.81 ± 0.01	0.71 ± 0.01	0.67 ± 0.05	4
Band-ignore CE	none	0.64 ± 0.07	0.65 ± 0.07	0.47 ± 0.19	0.48 ± 0.05	4
CE	boundary_p0.3	0.78 ± 0.04	0.81 ± 0.04	0.69 ± 0.03	0.65 ± 0.09	4
SCE	boundary_p0.3	0.70 ± 0.09	0.80 ± 0.03	0.65 ± 0.07	0.38 ± 0.27	4
GCE	boundary_p0.3	0.78 ± 0.02	0.81 ± 0.02	0.70 ± 0.02	0.64 ± 0.04	4
Band-ignore CE	boundary_p0.3	0.59 ± 0.08	0.63 ± 0.08	0.42 ± 0.22	0.36 ± 0.13	4
CE	flip_p0.3	0.53 ± 0.10	0.54 ± 0.09	0.28 ± 0.14	0.32 ± 0.19	4
SCE	flip_p0.3	0.53 ± 0.13	0.57 ± 0.12	0.36 ± 0.23	0.21 ± 0.25	4
GCE	flip_p0.3	0.59 ± 0.09	0.60 ± 0.07	0.43 ± 0.08	0.38 ± 0.25	4
Band-ignore CE	flip_p0.3	0.45 ± 0.09	0.46 ± 0.06	0.23 ± 0.15	0.17 ± 0.19	4
Band-ignore CE	flip_p0.1	0.52 ± 0.12	0.54 ± 0.10	0.31 ± 0.25	0.28 ± 0.15	3
Band-ignore CE	flip_p0.2	0.44 ± 0.09	0.46 ± 0.06	0.22 ± 0.19	0.14 ± 0.19	3




Fig. 3. Per-class test IoU of CE at $N = 500$ for clean, boundary-noise and flip-noise training labels.

class confusion.

H3 (flip damage is larger with less data): refuted / inconclusive, opposite direction. The drop caused by flips (clean minus flip mIoU, mean over four seeds) is 0.19 at $N = 100$ and 0.26 at $N = 2000$; the difference of drops, $N = 100$ minus $N = 2000$, is -0.06 (Welch $p = 0.25$). The registered direction is not supported: there is no evidence that more data dilutes object-flip noise, and the point estimate goes the other way. Over seeds 0–2 shared by all sizes the drop is 0.20, 0.22, 0.28, 0.28, 0.27 for $N = 100, 250, 500, 1000, 2000$, increasing at first and then flat. For boundary noise the drops are tiny at both sizes (0.011 and 0.001; difference test $p = 0.26$). With flips, the seed-to-seed std of CE grows with N (0.03 at $N = 100$, 0.11 at $N = 2000$): some seeds recover a decent triangle class and others do not, and four seeds cannot resolve this.

H4 (robust losses beat CE under noise and do not lose on clean data): not supported. Table II lists the registered

TABLE II

Registered comparisons against CE at $N = 500$ (mIoU; difference is CE minus system, so negative means the system is better; four seeds, Welch p).

Task	System	CE–system	Welch p
clean	GCE	0.003	0.888
clean	SCE	0.030	0.198
clean	Band-ignore CE	0.160	0.014
boundary 0.3	GCE	-0.001	0.983
boundary 0.3	SCE	0.082	0.166
boundary 0.3	Band-ignore CE	0.196	0.010
flip 0.3	GCE	-0.065	0.384
flip 0.3	SCE	0.002	0.982
flip 0.3	Band-ignore CE	0.080	0.294

comparisons (rh compare, CE as reference). GCE equals CE on clean and boundary-noise labels and has a higher mean under flips, but the effect is not significant (CE minus GCE -0.07 , $p = 0.38$); a seed-paired test, which we do not use for the verdict, gives $p = 0.019$, so the direction is consistent across seeds but its size is not established. SCE is not better than CE on any task and is worse on boundary noise in mean (CE minus SCE 0.08, $p = 0.17$), where its triangle IoU is 0.38 with a std of 0.27 over seeds. Our band-ignore CE is clearly worse than CE on clean labels (difference 0.16, $p = 0.014$) and on boundary noise (0.20, $p = 0.010$), i.e. the method designed for boundary noise fails exactly there; under flips it is lower in mean (0.08, $p = 0.29$). A plausible but untested reason is that, in a 64×64 image with small objects, a two-pixel band removes supervision on a large share of the foreground pixels and on the true edges that define square and triangle corners; we did not measure the ignored fraction.

Failure cases. (i) Band-ignore CE as above. (ii) Triangle IoU of GCE and CE under flips has a per-seed spread of several tenths (Table I), because in some seeds the triangle class is learned and in others it is largely predicted as

TABLE III

Training-set size: CE test IoU (mean $\pm$ std, n seeds), all tasks and sizes run.

Method	Task	miou $\uparrow$	iou_circle $\uparrow$	iou_square $\uparrow$	iou_triangle $\uparrow$	n
CE $n=100$	none	0.68 ± 0.04	0.64 ± 0.06	0.55 ± 0.05	0.56 ± 0.06	4
CE $n=250$	none	0.76 ± 0.03	0.77 ± 0.04	0.66 ± 0.04	0.65 ± 0.05	3
CE $n=2000$	none	0.82 ± 0.03	0.84 ± 0.02	0.74 ± 0.03	0.71 ± 0.05	4
CE $n=1000$	none	0.81 ± 0.03	0.83 ± 0.03	0.72 ± 0.03	0.70 ± 0.05	3
CE $n=100$	boundary_p0.3	0.67 ± 0.04	0.63 ± 0.05	0.54 ± 0.05	0.55 ± 0.05	4
CE $n=2000$	boundary_p0.3	0.82 ± 0.01	0.84 ± 0.02	0.73 ± 0.02	0.71 ± 0.02	4
CE $n=100$	flip_p0.3	0.49 ± 0.03	0.35 ± 0.07	0.30 ± 0.03	0.32 ± 0.03	4
CE $n=250$	flip_p0.3	0.54 ± 0.02	0.47 ± 0.04	0.31 ± 0.02	0.39 ± 0.03	3
CE $n=2000$	flip_p0.3	0.56 ± 0.11	0.54 ± 0.11	0.30 ± 0.11	0.41 ± 0.24	4
CE $n=1000$	flip_p0.3	0.53 ± 0.14	0.53 ± 0.11	0.23 ± 0.20	0.36 ± 0.27	3

TABLE IV

Noise level, CE, flip noise, $N = 500$ (three seeds; clean and $p = 0.3$ over the same three seeds).

Method	Task	miou $\uparrow$	iou_circle $\uparrow$	iou_square $\uparrow$	iou_triangle $\uparrow$	n
CE	flip_p0.1	0.69 ± 0.02	0.69 ± 0.03	0.53 ± 0.04	0.57 ± 0.02	3
CE	flip_p0.2	0.63 ± 0.04	0.62 ± 0.05	0.43 ± 0.06	0.50 ± 0.03	3

another class. (iii) SCE has the lowest clean-label mean of the three published-loss systems (CE minus SCE 0.03, $p = 0.20$); the reverse-CE term may make optimisation slower in our short budget, so the baseline may be disadvantaged, but we did not test this and did not retune its defaults.

VI. Ablations and Analysis

Four sensitivity studies accompany the main grid. Their groups are small (three or four seeds), so we describe patterns rather than test them.

Training-set size (Table III, Fig. 2). Under clean and boundary-noise labels the curve is the same; under flips mIoU is nearly flat in N (0.49 at $N = 100$, 0.56 at $N = 2000$), so the gap to clean labels widens with data (Sec. V, H3). As a direct memorisation indicator, we log how often the final CE network agrees with the noisy versus the clean training labels on its own training pixels under flips. At $N = 100$ it agrees more with the noisy labels (0.97 vs 0.90), i.e. it has fitted flipped objects; at $N = 500$ (0.88 vs 0.89) and $N = 2000$ (0.87 vs 0.90) it agrees slightly more with the clean labels. Because every size gets the same number of optimiser steps, small N means more epochs, which is a confound we did not separate from N itself.

Flip-noise level (Table IV). Over seeds 0–2, mean mIoU is 0.79 (clean), 0.69 ($p = 0.1$), 0.63 ($p = 0.2$) and 0.51 ($p = 0.3$), with stds of 0.02, 0.04 and 0.12 for the three noisy levels: the loss is smooth and monotone in the mean across the three levels run, but the spread grows sharply at $p = 0.3$. Only these levels were run, so we make no claim about other p , and no boundary-noise level sweep was run. The main group also has band-ignore CE at $p = 0.1, 0.2$ (Table I); there it again tracks lower than CE at the same level.

GCE exponent (Table V). mIoU is 0.54 at $q = 0.3$ and 0.45 at $q = 0.9$, against 0.59 for the default $q = 0.7$ (different seed set). The largest q is clearly worst in mean (triangle IoU 0.12, close to a collapse), consistent with

TABLE V

GCE exponent q under flip $p = 0.3$, $N = 500$ (three seeds). $q = 0.7$ is the four-seed main row.

Method	miou $\uparrow$	iou_circle $\uparrow$	iou_square $\uparrow$	iou_triangle $\uparrow$	n
GCE $q=0.3$	0.54 ± 0.11	0.51 ± 0.14	0.36 ± 0.08	0.32 ± 0.25	3
GCE $q=0.9$	<u>0.45 ± 0.15</u>	<u>0.48 ± 0.14</u>	<u>0.23 ± 0.25</u>	<u>0.12 ± 0.21</u>	3

TABLE VI

CE training length at $N = 100$ under flip $p = 0.3$ (three seeds; the 600-step row is in Table III).

Method	miou $\uparrow$	iou_circle $\uparrow$	iou_square $\uparrow$	iou_triangle $\uparrow$	n
CE steps=200	0.34 ± 0.01	0.23 ± 0.11	0.22 ± 0.09	0.00 ± 0.00	3
CE steps=1200	0.49 ± 0.02	0.30 ± 0.05	0.36 ± 0.03	0.32 ± 0.04	3

the known under-fitting of MAE-like losses; $q = 0.3$ is close to CE (0.53), as expected since GCE tends to CE as $q \rightarrow 0$. The main-table GCE therefore uses a default setting that was not chosen on this sweep, and the best q is not identified.

Training length (Table VI). At $N = 100$, 200 steps give mIoU 0.34 (triangle IoU 0.00: the class is not yet learned), 600 steps give 0.49, and 1200 steps give 0.49. So mIoU rises from 200 to 600 steps and is essentially unchanged from 600 to 1200 steps: we did not observe the late-training collapse that the early-learning literature predicts [2], [10]. The network may begin to memorise flips only later (the train-fit statistics above show it does fit them at $N = 100$), but the mIoU at 1200 steps does not show a loss.

VII. Limitations

Scale and seeds. Four seeds in the main grid and at $N = 100, 2000$, three in the sweeps and at $N = 250, 1000$, against five or more requested; with these n the Welch tests have little power, per-seed variance under flip noise is large (std up to 0.14 mIoU), and a non-significant result is not evidence of equality. No multiplicity correction was applied.

One architecture, one domain, one budget. A 120356-parameter U-Net on 64×64 synthetic shapes with a fixed 600-step budget for every N . We did not tune the optimiser per system or per noise level, so baselines such as SCE (and the large- q GCE) may be handicapped by the short budget; the optimiser budget was selected from pilot runs including two single-seed clean-label runs scored on the test split, a small leak that applies equally to all systems. A full study would use a held-out clean validation set for model selection, a longer schedule per N , real datasets with real annotation noise, and several architectures.

Noise models and levels. Boundary noise was run at one level and only at $N = 100, 500, 2000$; the two noise models are not matched in the number of corrupted pixels, so the flip-vs-boundary comparison (H2) compares nominal levels, not equal corruption. Flips are class-symmetric and object-wide, a simple model of real annotation errors. We reimplemented only GCE, SCE and our band-ignore loss;

no sample-selection, co-teaching or early-learning method and no Dice-type loss was evaluated, so the negative result for robust losses covers only the pixel-wise losses tested.

Confounds and untested mechanisms. Fixed steps means small N is trained for more epochs, confounding size with epochs; the explanation for band-ignore CE's failure and the claim that flip damage is a foreground-class confusion beyond the per-class IoU are not isolated by any run.

VIII. Conclusion

In a small, exactly-known synthetic setting, training-set size helped on clean labels (H1) and object-level class flips hurt far more than boundary jitter at the same nominal level (H2), with the damage concentrated in the square and triangle classes. The hoped-for dilution of flip noise by more data (H3) did not appear, and none of the pixel-wise robust losses we reimplemented was demonstrably better than plain cross-entropy (H4); the one that came closest, GCE, did not reach significance with four seeds, and our own band-ignore loss was worse than CE where it was meant to help. These conclusions are limited to this architecture, budget and noise model, and to four (or three) seeds per cell; longer-schedule, validation-selected and real-data follow-ups are needed before they would justify advice to practitioners.

References

- [1] C. Zhang, S. Bengio, M. Hardt, B. Recht, and O. Vinyals, "Understanding deep learning requires rethinking generalization," arXiv preprint arXiv:1611.03530, 2016. [Online]. Available: <https://arxiv.org/abs/1611.03530>
- [2] D. Arpit, S. Jastrzebski, N. Ballas, D. Krueger, E. Bengio, M. S. Kanwal, T. Maharaj, A. Fischer, A. Courville, Y. Bengio, and S. Lacoste-Julien, "A Closer Look at Memorization in Deep Networks," arXiv preprint arXiv:1706.05394, 2017. [Online]. Available: <https://arxiv.org/abs/1706.05394>
- [3] D. Rolnick, A. Veit, S. Belongie, and N. Shavit, "Deep Learning is Robust to Massive Label Noise," arXiv preprint arXiv:1705.10694, 2017. [Online]. Available: <https://arxiv.org/abs/1705.10694>
- [4] H. Song, M. Kim, D. Park, Y. Shin, and J.-G. Lee, "Learning from Noisy Labels with Deep Neural Networks: A Survey," arXiv preprint arXiv:2007.08199, 2020. [Online]. Available: <https://arxiv.org/abs/2007.08199>
- [5] O. Ronneberger, P. Fischer, and T. Brox, "U-Net: Convolutional Networks for Biomedical Image Segmentation," arXiv preprint arXiv:1505.04597, 2015. [Online]. Available: <https://arxiv.org/abs/1505.04597>
- [6] Z. Zhang and M. R. Sabuncu, "Generalized Cross Entropy Loss for Training Deep Neural Networks with Noisy Labels," arXiv preprint arXiv:1805.07836, 2018. [Online]. Available: <https://arxiv.org/abs/1805.07836>
- [7] Y. Wang, X. Ma, Z. Chen, Y. Luo, J. Yi, and J. Bailey, "Symmetric Cross Entropy for Robust Learning with Noisy Labels," arXiv preprint arXiv:1908.06112, 2019. [Online]. Available: <https://arxiv.org/abs/1908.06112>
- [8] F. Milletari, N. Navab, and S.-A. Ahmadi, "V-Net: Fully Convolutional Neural Networks for Volumetric Medical Image Segmentation," arXiv preprint arXiv:1606.04797, 2016. [Online]. Available: <https://arxiv.org/abs/1606.04797>
- [9] B. Han, Q. Yao, X. Yu, G. Niu, M. Xu, W. Hu, I. Tsang, and M. Sugiyama, "Co-teaching: Robust Training of Deep Neural Networks with Extremely Noisy Labels," arXiv preprint arXiv:1804.06872, 2018. [Online]. Available: <https://arxiv.org/abs/1804.06872>
- [10] S. Liu, J. Niles-Weed, N. Razavian, and C. Fernandez-Granda, "Early-Learning Regularization Prevents Memorization of Noisy Labels," arXiv preprint arXiv:2007.00151, 2020. [Online]. Available: <https://arxiv.org/abs/2007.00151>
- [11] D. Karimi, H. Dou, S. K. Warfield, and A. Gholipour, "Deep learning with noisy labels: exploring techniques and remedies in medical image analysis," arXiv preprint arXiv:1912.02911, 2019. [Online]. Available: <https://arxiv.org/abs/1912.02911>
- [12] J. Hestness, S. Narang, N. Ardalani, G. Diamos, H. Jun, H. Kianinejad, M. M. A. Patwary, Y. Yang, and Y. Zhou, "Deep Learning Scaling is Predictable, Empirically," arXiv preprint arXiv:1712.00409, 2017. [Online]. Available: <https://arxiv.org/abs/1712.00409>